

```
In [8]: import pandas as pd
import numpy as np
import tensorflow as tf
import tensorflow_hub as hub
from tensorflow.keras.models import Model
```

```
In [9]: reviews=pd.read_csv(':/content/drive/MyDrive/NLP/preprocessed.csv')
```

```
In [10]: reviews.head()
```

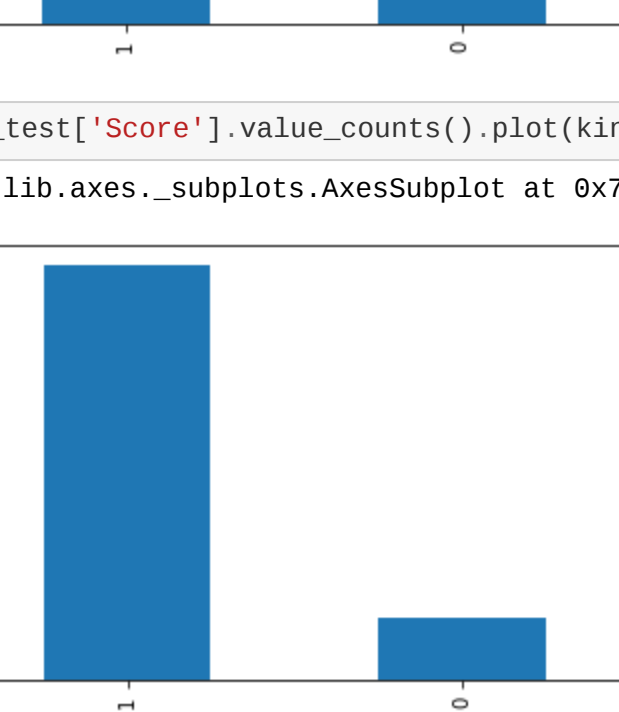
Out[10]:

	Text	Score	len
0	The tea was of great quality and it tasted lik...	1	30
1	My cat loves this. The pellets are nice and s...	1	31
2	Great product. Does not really get rid of...	1	41
3	This gum is my favorite! I would advise every...	1	27
4	I also found out about this product because of...	1	22

```
In [11]: from sklearn.model_selection import train_test_split
reviews_train, reviews_test = train_test_split(reviews, test_size=0.20, st
ratify = reviews['Score'], random_state=33)
```

```
In [12]: #plot bar graphs of train and test Scores
reviews_train['Score'].value_counts().plot(kind='bar')
```

Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x7f695d392c50>



```
In [13]: reviews_train['Score'].value_counts().plot(kind='bar')
```

Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f695d287c50>

```
In [14]: X_train, y_train = np.array(reviews_train['Text']), np.array(reviews_train
['Score'])
X_test, y_test = np.array(reviews_test['Text']), np.array(reviews_test['S
core'])
```

Part2 - Creating BERT Model

```
In [15]: # Loading the Pretrained Model from tensorflow HUB
tf.keras.backend.clear_session()

# maximum length of a seq in the data we have, for now i am making it as
55. You can change this
max_seq_length = 128 # maximum seq length should be less than 512
#BERT takes 3 inputs

#this is input words. Sequence of words represented as integers
input_word_ids = tf.keras.layers.Input(shape=(max_seq_length,), dtype=tf
.int32, name="input_word_ids")

#Embed vector if you are padding anything
input_mask = tf.keras.layers.Input(shape=(max_seq_length,), dtype=tf.int
32, name="input_mask")

#segment vectors. If you are giving only one sentence for the classifica
tion, total seq vector is 0.
#If you are giving two sentences with [sep] token separated, first seq s
egment vectors are zeros and
#second seq segment vector are 1's
segment_ids = tf.keras.layers.Input(shape=(max_seq_length,), dtype=tf.in
t32, name="segment_ids")

#bert layer
bert = tf.keras.layers.experimental.preprocessing.BertLayer(
    d_model=768, d_ff=4096, num_heads=12, num_encoder_layers=12, num_decoder_layers=12,
    dropout=0.1, trainable=True)

#BERT model
bert_model = Model(inputs=[input_word_ids, input_mask, segment_ids], out
puts=pooled_output)
```

```
In [16]: bert_model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #	Connect ed to
input_word_ids (InputLayer)	[(None, 128)]	0	[]
input_mask (InputLayer)	[(None, 128)]	0	[]
segment_ids (InputLayer)	[(None, 128)]	0	[]
keras_layer (KerasLayer)	[(None, 768),	109482241	['input
bert_model (BertModel)	(None, 128, 768)]		_mask[0][0]',
			'segment_ids[0][0]']

Total params: 109,482,241
Trainable params: 0
Non-trainable params: 109,482,241

```
In [17]: bert_model.output
```

Out[17]: <KerasTensor: shape=(None, 768) dtype=float32 (created by layer 'keras_layer')>

Part3: Tokenization

```
In [18]: #getting Vocab File
vocab_file = bert_layer.resolved_object.vocab_file.asset_path.numpy()
do_lower_case = bert_layer.resolved_object.do_lower_case.numpy()
```

```
In [20]: # train initial sentencepiece
Collecting sentencepiece
Downloading sentencepiece-0.1.96-cp37-cp37m-manylinux_2_17_x86_64.many
linux2014_x86_64.whl (1.2 MB)
Successfully installed sentencepiece-0.1.96
```

```
In [21]: import tokenization
```

```
In [22]: tokenizer = tokenization.FullTokenizer(vocab_file, do_lower_case)
```

Grader function 3

```
In [23]: #it has to give no error
def grader_tokenize(tokenizer):
    out = False
    try:
        out=('[CLS]' in tokenizer.vocab) and ('[SEP]' in tokenizer.vocab)
    except:
        out = False
    assert(out==True)
    return out
grader_tokenize(tokenizer)
```

Out[23]: True

```
In [24]: X_train_tokens = list(map(lambda t: ['[CLS]'] + tokenizer.tokenize(t) +
['[SEP]'], X_train))
X_test_tokens = list(map(lambda t: ['[CLS]'] + tokenizer.tokenize(t) +
['[SEP]'], X_test))
```

```
In [25]: for i in range(len(X_train_tokens)):
    offby = max_seq_length - len(X_train_tokens[i])
    if offby>0:
        X_train_tokens[i] = np.array(X_train_tokens[i] + ['[PAD]']*offby)
    else:
        X_train_tokens[i] = np.array(X_train_tokens[i][:max_seq_length-1] +
['[SEP]'])
```

```
In [26]: for i in range(len(X_test_tokens)):
    offby = max_seq_length - len(X_test_tokens[i])
    if offby>0:
        X_test_tokens[i] = np.array(X_test_tokens[i] + ['[PAD]']*offby)
    else:
        X_test_tokens[i] = np.array(X_test_tokens[i][:max_seq_length-1] +
['[SEP]'])
```

```
In [27]: train_tokens = np.array(list(map(tokenizer.convert_tokens_to_ids, X_train_tokens)))
test_tokens = np.array(list(map(tokenizer.convert_tokens_to_ids, X_test_tokens)))
```

```
In [28]: train_masks = np.zeros((train_tokens.shape[0], max_seq_length))
test_masks = np.zeros((test_tokens.shape[0], max_seq_length))
for j in range(len(train_tokens)):
    for i in range(len(train_tokens[j])):
        if train_tokens[j][i]!=0:
            train_masks[j][i]=1
for m in range(len(test_tokens)):
    for i in range(len(test_tokens[m])):
        if test_tokens[m][i]!=0:
            test_masks[m][i]=1
```

```
In [29]: train_segments = np.zeros((train_tokens.shape[0], train_tokens.shape[1]))
test_segments = np.zeros((test_tokens.shape[0], test_tokens.shape[1]))
```

```
In [30]: import pickle
pickle.dump(X_train, train_tokens, train_masks, train_segments, y_train,
open('/content/drive/MyDrive/NLP/train_data.pkl', 'wb'))
pickle.dump(X_test, test_tokens, test_masks, test_segments, y_test,
open('/content/drive/MyDrive/NLP/test_data.pkl', 'wb'))
```

```
In [31]: import pickle
X_train, X_train_tokens, X_train_mask, X_train_segment, y_train = pickle
.load(open('/content/drive/MyDrive/NLP/train_data.pkl', 'rb'))
X_test, X_test_tokens, X_test_mask, X_test_segment, y_test = pickle.load
(open('/content/drive/MyDrive/NLP/test_data.pkl', 'rb'))
```

Grader Function 4

```
In [32]: def grader_alltokens_train():
    out = False
    if type(X_train_tokens) == np.ndarray:
        temp_shapes = (X_train_tokens.shape[1]==max_seq_length) and (X_train_mask.shape[1]==max_seq_length) and \
(X_train_segment.shape[1]==max_seq_length)
        segment_temp = not np.any(X_train_segment)
        mask_temp = np.sum(X_train_mask==0) == np.sum(X_train_tokens==0)
        no_cls = np.sum(X_train_tokens==tokenizer.vocab['[CLS]'])==X_train_tokens.shape[0]
        no_sep = np.sum(X_train_tokens==tokenizer.vocab['[SEP]'])==X_train_tokens.shape[0]
        out = temp_shapes and segment_temp and mask_temp and no_cls and no_sep
    else:
        print('Type of all above token arrays should be numpy array not list')
    out = False
    assert(out==True)
    return out
grader_alltokens_train()
```

Out[32]: True

Grader Function 5

```
In [33]: def grader_alltokens_test():
    out = False
    if type(X_test_tokens) == np.ndarray:
        temp_shapes = (X_test_tokens.shape[1]==max_seq_length) and (X_test_mask.shape[1]==max_seq_length) and \
(X_test_segment.shape[1]==max_seq_length)
        segment_temp = not np.any(X_test_segment)
        mask_temp = np.sum(X_test_mask==0) == np.sum(X_test_tokens==0)
        no_cls = np.sum(X_test_tokens==tokenizer.vocab['[CLS]'])==X_test_tokens.shape[0]
        no_sep = np.sum(X_test_tokens==tokenizer.vocab['[SEP]'])==X_test_tokens.shape[0]
        out = temp_shapes and segment_temp and mask_temp and no_cls and no_sep
    else:
        print('Type of all above token arrays should be numpy array not list')
    out = False
    assert(out==True)
    return out
grader_alltokens_test()
```

Out[33]: True

Part 4 - Getting Embeddings from BERT Model

```
In [34]: bert_model.input
<KerasTensor: shape=(None, 128) dtype=int32 (created by layer 'input_word_ids')>
<KerasTensor: shape=(None, 128) dtype=int32 (created by layer 'input_mask')>
<KerasTensor: shape=(None, 128) dtype=int32 (created by layer 'segment_ids')>
```

```
In [35]: bert_model.output
```

Out[35]: <KerasTensor: shape=(None, 768) dtype=float32 (created by layer 'keras_layer')>

```
In [36]: X_train_pooled_output=bert_model.predict([X_train_tokens,X_train_mask,X_train_segment])
```

```
In [37]: X_test_pooled_output=bert_model.predict([X_test_tokens,X_test_mask,X_test_segment])
```

```
In [38]: pickle.dump(X_train_pooled_output, X_test_pooled_output), open('/content/drive/MyDrive/NLP/train_output.pkl', 'wb'))
```

```
In [39]: X_train_pooled_output, X_test_pooled_output= pickle.load(open('/content/drive/MyDrive/NLP/train_output.pkl', 'rb'))
```

Grader Function 6

```
In [40]: #now we have X_train_pooled_output, y_train
#X_test_pooled_output, y_test
def grader_output():
    assert(X_train_pooled_output.shape[1]==768)
    assert(len(y_train)==len(X_train_pooled_output))
    assert(X_test_pooled_output.shape[1]==768)
    assert(len(y_test)==len(X_test_pooled_output))
    assert(len(y_train.shape)==1)
    assert(len(X_train_pooled_output.shape)==2)
    assert(len(y_test.shape)==1)
    assert(len(X_test_pooled_output.shape)==2)
    return True
grader_output()
```

Out[40]: True

Part 5 - Training a NN with 768 features

```
In [41]: from sklearn.metrics import roc_auc_score
def auc(y_true, y_pred):
    return roc_auc_score(y_true, y_pred, average='micro')
```

```
In [42]: from tensorflow.keras.layers import Input, Dense, Activation, Dropout, LSTM
from tensorflow.keras.models import Model
```

```
In [43]: import tensorflow as tf
y_train = tf.keras.utils.to_categorical(y_train, 2)
y_test = tf.keras.utils.to_categorical(y_test, 2)
```

```
In [52]: initializer = tf.keras.initializers.he_normal()
input_layer = Input(shape=(X_train_pooled_output.shape[1],), name='input_data')
layer1 = Dense(20, activation='relu', kernel_initializer=initializer, name = 'Dense_1')(input_layer)
layer2 = Dense(15, activation='relu', kernel_initializer=initializer, name = 'Dense_2')(layer1)
layer3 = Dense(10, activation='relu', kernel_initializer=initializer, name = 'Dense_3')(layer2)
layer4 = Dense(5, activation='relu', kernel_initializer=initializer, name = 'Dense_4')(layer3)
output = Dense(2, activation='softmax', kernel_initializer=initializer, name = 'output_layer_to_classify_with_softmax')(layer4)
```

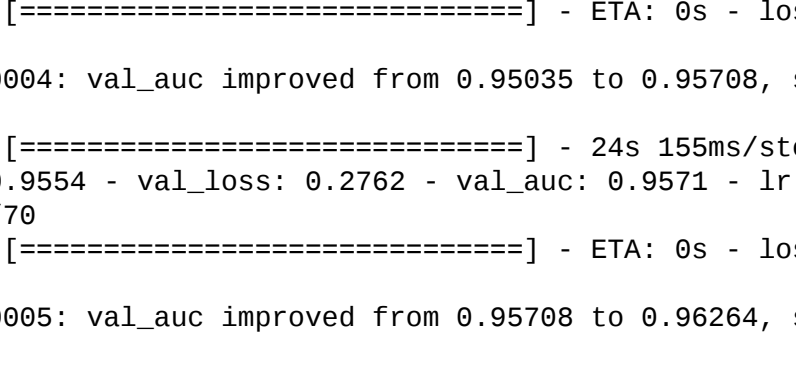
```
model = Model(inputs=input_layer, outputs=output)
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
loss='categorical_crossentropy', metrics=[auc], run_eagerly=True)
model.summary()
```

Model: "model_3"

Layer (type)	Output Shape	Param #
input_data (InputLayer)	[(None, 768)]	0
Dense_1 (Dense)	(None, 20)	15380
Dense_2 (Dense)	(None, 15)	315
Dense_3 (Dense)	(None, 10)	160
Dense_4 (Dense)	(None, 5)	52
output_layer_to_classify_with_softmax (Dense)	(None, 2)	10

Total params: 15,922
Trainable params: 15,922
Non-trainable params: 0

```
In [53]: from tensorflow.keras.utils import plot_model
plot_model(model, to_file='NN for NLP Transfer Learning.png', show_shape=True)
```



```
In [54]: from tensorflow.keras.callbacks import TensorBoard, ModelCheckpoint, EarlyStopping, ReduceLROnPlateau
import datetime
log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = TensorBoard(log_dir=log_dir, histogram_freq=1)
es = EarlyStopping(monitor='val_auc', mode='max', verbose=1, patience=5)
mc = ModelCheckpoint('NLP_h5', monitor='val_auc', mode='max', save_best_only=True, verbose=1)
reduce_lr = ReduceLROnPlateau(monitor='val_auc', factor=0.8, patience=2, verbose=0)
callback_list = [tensorboard_callback, es, mc, reduce_lr]
```

```
In [55]: %load_ext tensorboard
The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard
```

```
In [56]: !rm -rf ./logs/
```

```
In [57]: model.fit(X_train_pooled_output, y_train, epochs=70, batch_size=512,
validation_data=(X_test_pooled_output, test), callbacks = callback_list)
```

```
Epoch 1/70
157/157 [=====] - ETA: 0s - loss: 0.3900 - auc: 0.8890
Epoch 00001: val_auc improved from -inf to 0.91388, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.3900 - auc: 0.8890
Epoch 2/70
157/157 [=====] - ETA: 0s - loss: 0.3534 - auc: 0.9291
Epoch 00002: val_auc improved from 0.91188 to 0.93778, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.3534 - auc: 0.9291
Epoch 3/70
157/157 [=====] - ETA: 0s - loss: 0.3136 - auc: 0.9467
Epoch 00003: val_auc improved from 0.93778 to 0.95035, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.3136 - auc: 0.9467
Epoch 4/70
157/157 [=====] - ETA: 0s - loss: 0.2847 - auc: 0.9554
Epoch 00004: val_auc improved from 0.95035 to 0.95708, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.2847 - auc: 0.9554
Epoch 5/70
157/157 [=====] - ETA: 0s - loss: 0.2657 - auc: 0.9607
Epoch 00005: val_auc improved from 0.95708 to 0.96264, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.2657 - auc: 0.9607
Epoch 6/70
157/157 [=====] - ETA: 0s - loss: 0.2505 - auc: 0.9648
Epoch 00006: val_auc improved from 0.96264 to 0.96515, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.2505 - auc: 0.9648
Epoch 7/70
157/157 [=====] - ETA: 0s - loss: 0.2387 - auc: 0.9677
Epoch 00007: val_auc improved from 0.96515 to 0.96743, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.2387 - auc: 0.9677
Epoch 8/70
157/157 [=====] - ETA: 0s - loss: 0.2294 - auc: 0.9697
Epoch 00008: val_auc improved from 0.96743 to 0.96992, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.2294 - auc: 0.9697
Epoch 9/70
157/157 [=====] - ETA: 0s - loss: 0.2231 - auc: 0.9715
Epoch 00009: val_auc did not improve from 0.96992
157/157 [=====] - ETA: 0s - loss: 0.2231 - auc: 0.9715
Epoch 10/70
157/157 [=====] - ETA: 0s - loss: 0.2176 - auc: 0.9728
Epoch 00010: val_auc improved from 0.96992 to 0.97185, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.2176 - auc: 0.9728
Epoch 11/70
157/157 [=====] - ETA: 0s - loss: 0.2132 - auc: 0.9734
Epoch 00011: val_auc improved from 0.97185 to 0.97338, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.2132 - auc: 0.9734
Epoch 12/70
157/157 [=====] - ETA: 0s - loss: 0.2100 - auc: 0.9741
Epoch 00012: val_auc improved from 0.97338 to 0.97421, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.2100 - auc: 0.9741
Epoch 13/70
157/157 [=====] - ETA: 0s - loss: 0.2083 - auc: 0.9748
Epoch 00013: val_auc improved from 0.97421 to 0.97438, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.2083 - auc: 0.9748
Epoch 14/70
157/157 [=====] - ETA: 0s - loss: 0.2073 - auc: 0.9754
Epoch 00014: val_auc improved from 0.97438 to 0.97455, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.2073 - auc: 0.9754
Epoch 15/70
157/157 [=====] - ETA: 0s - loss: 0.2060 - auc: 0.9764
Epoch 00015: val_auc did not improve from 0.97455
157/157 [=====] - ETA: 0s - loss: 0.2060 - auc: 0.9764
Epoch 16/70
157/157 [=====] - ETA: 0s - loss: 0.2000 - auc: 0.9768
Epoch 00016: val_auc did not improve from 0.97455
157/157 [=====] - ETA: 0s - loss: 0.2000 - auc: 0.9768
Epoch 17/70
157/157 [=====] - ETA: 0s - loss: 0.1988 - auc: 0.9768
Epoch 00017: val_auc improved from 0.97455 to 0.97571, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.1988 - auc: 0.9768
Epoch 18/70
157/157 [=====] - ETA: 0s - loss: 0.1978 - auc: 0.9768
Epoch 00018: val_auc improved from 0.97571 to 0.97596, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.1978 - auc: 0.9768
Epoch 19/70
157/157 [=====] - ETA: 0s - loss: 0.1963 - auc: 0.9771
Epoch 00019: val_auc improved from 0.97596 to 0.97620, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.1963 - auc: 0.9771
Epoch 20/70
157/157 [=====] - ETA: 0s - loss: 0.1956 - auc: 0.9772
Epoch 00020: val_auc improved from 0.97620 to 0.97627, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.1956 - auc: 0.9772
Epoch 21/70
157/157 [=====] - ETA: 0s - loss: 0.1935 - auc: 0.9772
Epoch 00021: val_auc improved from 0.97627 to 0.97627, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.1935 - auc: 0.9772
Epoch 22/70
157/157 [=====] - ETA: 0s - loss: 0.1925 - auc: 0.9772
Epoch 00022: val_auc improved from 0.97627 to 0.97627, saving model to NLP_h5
157/157 [=====] - ETA: 0s - loss: 0.1925 - auc: 0.9772
Epoch 23/70
157/157 [=====] - ETA: 0s - loss: 0.1922 - auc: 0.9772
Epoch 00023: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1922 - auc: 0.9772
Epoch 24/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00024: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 25/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00025: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 26/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00026: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 27/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00027: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 28/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00028: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 29/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00029: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 30/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00030: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 31/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00031: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 32/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00032: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 33/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00033: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 34/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00034: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 35/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00035: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 36/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00036: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 37/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00037: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 38/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00038: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 39/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00039: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 40/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00040: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 41/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00041: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 42/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00042: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 43/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00043: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 44/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00044: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 45/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00045: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 46/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00046: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 47/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00047: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 48/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00048: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 49/70
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
Epoch 00049: val_auc did not improve from 0.97689
157/157 [=====] - ETA: 0s - loss: 0.1921 - auc: 0.9772
```

```
Out[57]: <keras.callbacks.History at 0x7f68e2e45090>
```

```
In [58]: %tensorboard --logdir logs/fit
```

```
In [59]: ! unzip /content/test.csv-20211125113253-001.zip
Archive: /content/test.csv-20211125113253-001.zip
inflating: test.csv
```



```
In [60]: test_df = pd.read_csv('/content/test.csv')

In [61]: test_df.head()

Out[61]:
```

Text

0

Just opened Greenies Joint Care (individually ...

1

This product rocks :) My moon was very happy w...

2

The product was fine, but the cost of shipping...

3

I love this soup, it's great as part of a meal...

4

Getting ready to order again. These are great ...

```
In [62]: import re as re
def remove_tags(string):
    result = re.sub('<.*?>', '', string)
    return result

In [63]: def transfer(df, model, max_seq_length, bert_model):
df['Text'] = df['Text'].apply(lambda cw : remove_tags(cw))
test = np.array(df['Text'])
pre_token = list(map(lambda t: ['[CLS]'] + tokenizer.tokenize(t) + ['[SEP]'], test))
for i in range(len(pre_token)):
    offby = max_seq_length - len(pre_token[i])
    if offby>0:
        pre_token[i] = np.array(pre_token[i] + ['[PAD]'] * offby)
    else:
        pre_token[i] = np.array(pre_token[i][:max_seq_length-1] + ['[SEP]'])
tokens_array = np.array(list(map(tokenizer.convert_tokens_to_ids, pre_token)))
masks_array = np.zeros((tokens_array.shape[0], max_seq_length))
for k in range(len(tokens_array)):
    for j in range(len(tokens_array[k])):
        if tokens_array[k][j]!=0:
            masks_array[k][j] = 1
segments_array = np.zeros((tokens_array.shape[0], tokens_array.shape[1]))
pooled_output = bert_model.predict([tokens_array, masks_array, segments_array])
predictions = model.predict(pooled_output)
# Generate arg maxes for predictions
classes = np.argmax(predictions, axis = 1)
unique, counts = np.unique(classes, return_counts=True)
return dict(zip(unique, counts))

In [64]: # count of datapoints classified as 1 or 0 in the form of dictionary
out = transfer(test_df, model, max_seq_length, bert_model)
print(out)

{0: 36, 1: 316}
```

Observations:

1. Out of 352 datapoints in the test dataset, model classified 36 points as 0 and 316 datapoints as 1.

The following are the steps performed while solving this assignment:

2. Preprocessing the reviews csv file by removing all the html tags, assigning the score values as 1 if score is >3 and score value as 0 if score is <=2. Also dropping all the rows if the score value is 3.

3. Splitting into train and test data which gives X_train, X_test, y_train and y_test and checking if train and test data have the same distribution.

4. Creating tokens, masked array and segment arrays for X_train and X_test.

5. Getting Embeddings from BERT Model and extracting pooled_output, out of two outputs pooled_output, sequence_output which it produces.

6. Training the Neural Network using the embeddings obtained using BERT model without overfitting and underfitting.

7. Then, predicting the output for the test data which has not seen by the model.

8. Extracting the number of occurrences of classes 0 or 1.